



# Text Preprocessing in Natural Language Processing

Transforming Raw Text into Actionable Data

# Outline

- What is text preprocessing?
- Why is it important?
- Preprocessing techniques

Quick recap of previous lecture

# Why text preprocessing is important?

- **Noise Reduction:** Remove irrelevant characters (punctuation, HTML tags).
- **Consistency:** Lowercasing, standardizing formats (e.g., dates).
- **Efficiency:** Smaller vocabulary size = faster model training.
- **Accuracy:** Improves NLP task performance (e.g., sentiment analysis).

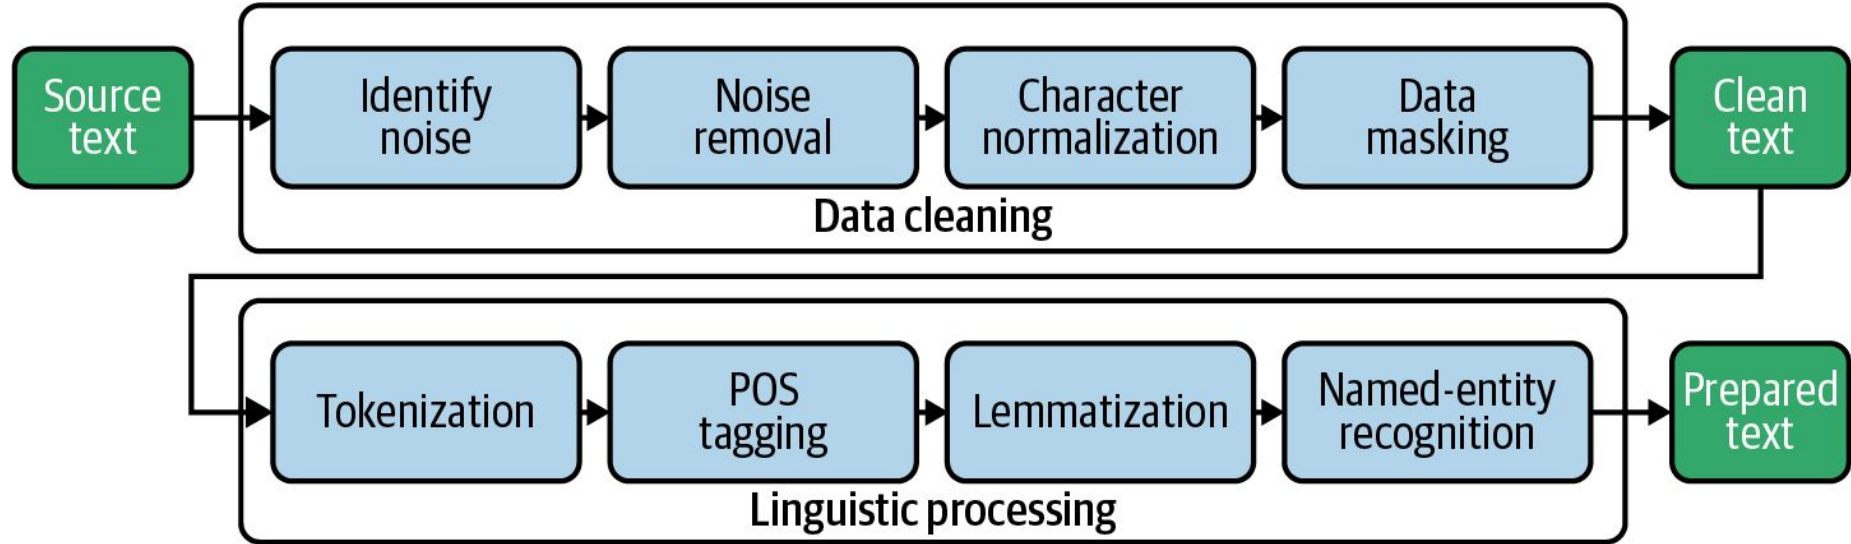
## Challenges

- **Ambiguity:** “Apple” (fruit vs. company).
- **Language Differences:** Morphology in Arabic vs. English.
- **Resource Limits:** Lemmatization requires heavy dictionaries.


# Domain-Specific Preprocessing

- **Social Media:** Emoji handling, slang normalization.
- **Scientific Texts:** Retain equations/symbols.
- **Medical Texts:** Protect sensitive terms (e.g., patient names).

# Text preprocessing stages



# Common preprocessing steps:



# Tokenization

Tokenization is the process of breaking down text into smaller chunks such as words or subwords.

## Importance

Foundation for most NLP tasks and models

**Tools:** NLTK and spaCy toolkits.

### Word tokenization

- Splits text into individual words. For example:
- Input: "Tokenization is fun!"
- Output: ["Tokenization", "is", "fun", "!"]

### Whitespace tokenization

- Splits text based on whitespace. For example:
- Input: "Tokenization is fun!"
- Output: ["Tokenization", "is", "fun!"]

### Subword tokenization

- Breaks words into smaller subword units
- Input: "unbelievable"
- Output: ["un", "believ", "able"]

### Character tokenization

- Splits text into individual characters. For example:
- Input: "Token"
- Output: ["T", "o", "k", "e", "n"]

### Sentence segmentation

- Splits text into sentences. For example:
- Input: "Tokenization is fun. Let's learn more!"
- Output: ["Tokenization is fun.", "Let's learn more!"]

### Regex-based tokenization

- Uses regular expressions to define custom tokenization rules.
- For example, splitting text based on punctuation or specific patterns.



# Tokenization

Tokenization is the process of breaking down text into smaller chunks such as words or subwords.

## Importance

Foundation for most NLP tasks and models

**Tools:** NLTK and spaCy toolkits..

## Sentence: Segmentation and Word/Subword: Tokenization

- **Subword Tokenization:** Handle rare/compound words (e.g., BPE in GPT).

### Word tokenization

- Splits text into individual words. For example:
- Input: "Tokenization is fun!"
- Output: ["Tokenization", "is", "fun", "!"]

### Whitespace tokenization

- Splits text based on whitespace. For example:
- Input: "Tokenization is fun!"
- Output: ["Tokenization", "is", "fun!"]

### Subword tokenization

- Breaks words into smaller subword units
- Input: "unbelievable"
- Output: ["un", "believ", "able"]

### Character tokenization

- Splits text into individual characters. For example:
- Input: "Token"
- Output: ["T", "o", "k", "e", "n"]

### Sentence segmentation

- Splits text into sentences. For example:
- Input: "Tokenization is fun. Let's learn more!"
- Output: ["Tokenization is fun.", "Let's learn more!"]

### Regex-based tokenization

- Uses regular expressions to define custom tokenization rules.
- For example, splitting text based on punctuation or specific patterns.

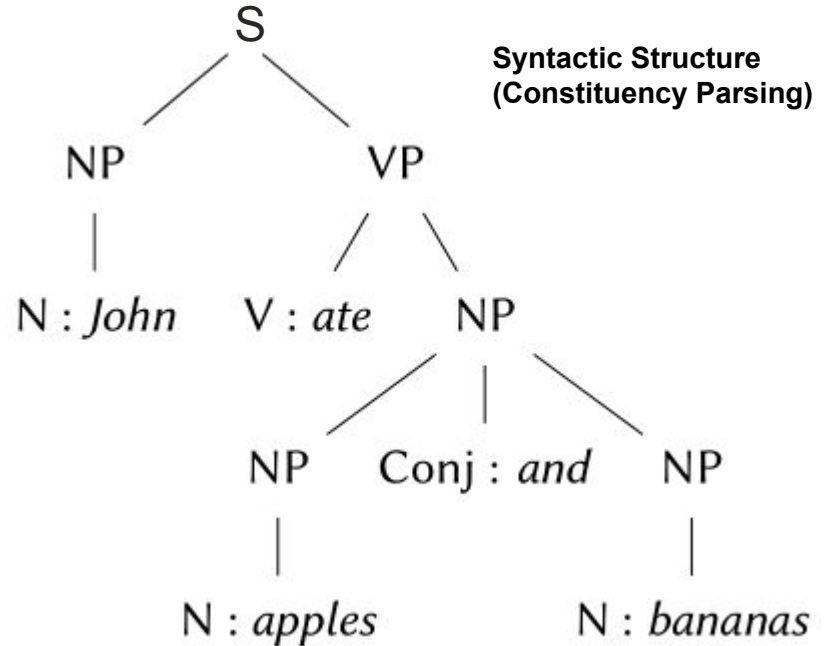

# Classes of tokenization algorithms

- Top-down tokenization
  - Rule-based (Penn Treebank Tokenizer)
- Bottom-up tokenization
  - Data-driven (Byte-Pair Encoding)
- Speed matters:
  - Tokenization is performed on every input before other language processing steps.
  - Large datasets and real-time systems need tokenizers that are highly optimized for speed.
  - Word tokenizers generally use deterministic (rule-based or algorithmic) logic for reliability and efficiency

# Top-down Tokenization

- Involves splitting text based on predefined rules or patterns.
- Focuses on higher-level linguistic units, typically words and sentences.
- Penn Treebank Tokenizer

## Dependency Parsing



# Top-down Tokenization

- Involves splitting text based on predefined rules or patterns.
- Focuses on higher-level linguistic units, typically words and sentences.
- Penn Treebank Tokenizer
  - Splits most punctuation from words.
    - Handles contractions and possessives (ownership)
    - e.g., "children's" → "children" + "'s"; "can't" → "ca" + "n't".

**Input:** "The San Francisco-based restaurant," they said,  
"doesn't charge \$10".

**Output:** "\_The\_San\_Francisco-based\_restaurant\_,\_"\_they\_said\_,\_  
"\_does\_n't\_charge\_\$\_10\_"\_.

# Top-down Tokenization

- Involves splitting text based on predefined rules or patterns.
- Focuses on higher-level linguistic units, typically words and sentences.
- Penn Treebank Tokenizer
  - Splits most punctuation from words.
    - Handles contractions and possessives (ownership)
    - e.g., "children's" → "children" + "'s"; "can't" → "ca" + "n't".
  - Advantages:
    - Fast and deterministic.
    - Designed for structured, well-formed text (e.g., newswire).
  - Deterministic output: Always produces the same tokenization for the same input

# Bottom-up Tokenization

- Constructs tokens from lower-level units (characters or bytes), merging them based on data frequencies.
- This data-driven approach is especially useful for handling **rare words** and **diverse language** inputs. [**Finite vocabulary from training data**]
  - NLP algorithm learn some facts from one corpus (**training corpus**) and make decisions on unseen corpus (**test corpus**)
  - By splitting words into smaller, reusable pieces (**subwords**), the model can generalize.
- Most tokenization schemes have two parts: **a token learner**, and **a token segmenter**.
  - Token learner learns the subword vocabulary from the training corpus (e.g., BPE training process).
  - Token segmenter applies the learned vocabulary to segment new (test) data into tokens (possibly splitting unknown words into familiar subwords).

# Subword Tokenization (Byte-Pair Encoding, BPE)

- Tokenizers automatically learn a set of tokens smaller than words (called **subwords**), enabling flexible splitting of both known and unknown words.

## Byte-Pair Encoding (BPE)

- Process:
  - Start with all characters as initial tokens.
  - Iteratively merge the most frequent adjacent token pairs in the corpus.
  - Build new tokens (subwords/words) with every merge step until  $k$  merges has been done.
    - $k = vocab\_size - initial\_vocab\_size$
- Outcome:
  - Handles arbitrary and out-of-vocabulary words.
  - Produces a compact, efficient vocabulary for downstream models
- Used by: Modern language models like OpenAI's GPT, Llama.

## corpus

5 l o w \_  
2 l o w e s t \_  
6 n e w e r \_  
3 w i d e r \_  
2 n e w \_

## vocabulary

\_, d, e, i, l, n, o, r, s, t, w

<i>Vocabulary</i>	<i>#count</i>
e	19
_	18
w	15
r	9
n	8
l	7
o	7
d	3
i	3
s	2
t	2

<i>Vocabulary</i>	<i>#count</i>
er	9
r_	9



## corpus

5 l o w \_  
2 l o w e s t \_  
6 n e w e r \_  
3 w i d e r \_  
2 n e w \_

## vocabulary

\_, d, e, i, l, n, o, r, s, t, w, er

<i>Vocabulary</i>	<i>#count</i>
e	19
_	18
w	15
r	9
n	8
l	7
o	7
d	3
i	3
s	2
t	2
er	9

<i>Vocabulary</i>	<i>#count</i>
er_	9
ne	9
lo	7
w_	7
ow	7
wer	6
wi	3
der	3
id	3
t_	2
we	2
st	2
t_	2

## corpus

5 l o w \_  
2 l o w e s t \_  
6 n e w e r\_  
3 w i d e r\_  
2 n e w \_

## vocabulary

\_, d, e, i, l, n, o, r, s, t, w, er, er\_

Vocabulary	#count
e	19
_	18
w	15
r	9
n	8
l	7
o	7
d	3
i	3
s	2
t	2
er	9
er_	9

Vocabulary	#count
ne	8
w_	7
lo	7
ow	7
wer_	6
der_	3
id	3
we	2
st	2
t_	2

**corpus**

5 l o w \_  
2 l o w e s t \_  
6 n e w er\_  
3 w i d er\_  
2 n e w \_

**vocabulary**

\_, d, e, i, l, n, o, r, s, t, w, er, er\_

**corpus**

5 l o w \_  
2 l o w e s t \_  
6 ne w er\_  
3 w i d er\_  
2 ne w \_

**vocabulary**

\_, d, e, i, l, n, o, r, s, t, w, er, er\_, ne

**corpus**

5 l o w \_  
 2 l o w e s t \_  
 6 ne w er\_  
 3 w i d er\_  
 2 ne w \_

**vocabulary**

\_, d, e, i, l, n, o, r, s, t, w, er, er\_, ne

**merge****current vocabulary**

(ne, w)	_, d, e, i, l, n, o, r, s, t, w, er, er_, ne, new
(l, o)	_, d, e, i, l, n, o, r, s, t, w, er, er_, ne, new, lo
(lo, w)	_, d, e, i, l, n, o, r, s, t, w, er, er_, ne, new, lo, low
(new, er_)	_, d, e, i, l, n, o, r, s, t, w, er, er_, ne, new, lo, low, newer_
(low, _)	_, d, e, i, l, n, o, r, s, t, w, er, er_, ne, new, lo, low, newer_, low_

```
function BYTE-PAIR ENCODING(strings  $C$ , number of merges  $k$ ) returns vocab  $V$ 

 $V \leftarrow$  all unique characters in  $C$            # initial set of tokens is characters
for  $i = 1$  to  $k$  do                           # merge tokens  $k$  times
     $t_L, t_R \leftarrow$  Most frequent pair of adjacent tokens in  $C$ 
     $t_{NEW} \leftarrow t_L + t_R$                  # make new token by concatenating
     $V \leftarrow V + t_{NEW}$                        # update the vocabulary
    Replace each occurrence of  $t_L, t_R$  in  $C$  with  $t_{NEW}$     # and update the corpus
return  $V$ 
```

**Figure 2.13** The token learner part of the BPE algorithm for taking a corpus broken up into individual characters or bytes, and learning a vocabulary by iteratively merging tokens. Figure adapted from [Bostrom and Durrett \(2020\)](#).

# Wordpiece Tokenizer

**function** BYTE-PAIR ENCODING(strings  $C$ , number of merges  $k$ ) **returns** vocab  $V$

$V \leftarrow$  all unique characters in  $C$       # initial set of tokens is characters

**for**  $i = 1$  **to**  $k$  **do**      # merge tokens  $k$  times

$t_L, t_R \leftarrow$  Most frequent pair of adjacent tokens in  $C$

$t_{NEW} \leftarrow t_L + t_R$       # make new token by concatenating

$V \leftarrow V + t_{NEW}$       # update the vocabulary

Replace each occurrence of  $t_L, t_R$  in  $C$  with  $t_{NEW}$       # and update the corpus

**return**  $V$

→  $\text{score} = (\text{freq\_of\_pair}) / (\text{freq\_of\_first\_element} \times \text{freq\_of\_second\_element})$

# Exercise for you:

Token learner: Find possible subwords using Wordpiece algorithm

**corpus**

5 l o w \_  
2 l o w e s t \_  
6 n e w e r \_  
3 w i d e r \_  
2 n e w \_

**vocabulary**

\_, d, e, i, l, n, o, r, s, t, w

$$\text{score} = (\text{freq\_of\_pair}) / (\text{freq\_of\_first\_element} \times \text{freq\_of\_second\_element})$$

# Word normalization

Word normalization is the task of putting words or tokens in a standard format.

- Case folding [generalize text to same case (lower case)]
- Normalize [(USA, US), (uh-huh, uhhuh)]
- Morphological analysis [mapping to same root word - (Warsaw, Warszawa)]
  - Stems
  - Affixes
    - Prefix
    - Suffix

Eg.

- cats - ['cat', 's']
- fox - ['fox']



# Lemmatization

Use linguistics to get valid root forms (e.g., “better” → “good”).

- Maps words to dictionary roots
- Produces valid words
- Requires POS tagging
- **The algorithm can be complex.**

was	→	(to) be
better	→	good
meeting	→	meeting

# Stemming

Reduce words to root form (e.g., “running” → “run”).

- Removes word suffixes aggressively
- May produce non-words
- Fast and simple

adjustable → adjust|  
formality → formalit|  
formality → form|al  
airliner → airlin|

# Text Normalization: Handling Contractions



Convert text to canonical form

Eg. Helloooo → Hello



Expand contractions

Convert shortened forms to full words



Standardize forms

Reduce variations in expression for consistency

**Tools:** Regex, custom dictionaries.

- can't
- cannot
- won't
- will not
- isn't
- is not

# Text Normalization: Emojis, Emoticons, and Spell Checking

## Handle emojis & emoticons

Convert to text descriptions or remove as needed

Toolkit: `demoji`

## Spell checking

Correct typos to improve data quality

**Use Case:** Fix typos in social media/text messages.

Toolkit: `pyspellchecker`



Afraid



Angry



Annoyed



Astonished



Bored



Confused



Content



Ecstatic



Gloomy



Happy



Miserable



Pleased



Sad



Satisfied



Serene



Surprised

# Minimum edit distance

Minimum Edit Distance is the **minimum number of operations** required to convert one string into another.

## Edit Operations:

- **Insertion** (add a character)
- **Deletion** (remove a character)
- **Substitution** (replace one character with another)

## Applications:

- Spell checking
- Machine translation evaluation
- Plagiarism detection
- Speech recognition

## Why does it matter?:

- MED quantifies how “similar” two pieces of text are [critical in many NLP tasks].

The gap between **intention** and **execution**

I	N	T	E	*	N	T	I	O	N
*	E	X	E	C	U	T	I	O	N
d	s	s				i	s		

# Minimum edit distance

Minimum Edit Distance is the **minimum number of operations** required to convert one string into another.

## Edit Operations:

- **Insertion** (add a character) [("cat" → "cart" by inserting 'r')]
- **Deletion** (remove a character) [("cart" → "cat" by deleting 'r')]
- **Substitution** (replace one character with another) [("cat" → "cut" by substituting 'a' with 'u')]

## Applications:

- Spell checking [("recieve" → "receive") distance = 2]
- Machine translation evaluation [Compare translated output with a reference translation]
- Plagiarism detection [Find passages that are almost the same but with small edits.]
- Speech recognition [Align output text with the actual spoken words to score accuracy.]

## Why does it matter?:

- MED quantifies how “similar” two pieces of text are [critical in many NLP tasks].

# How to calculate MED?

- The space of all possible edits is enormous, so we can't search naively.
- We could just remember the shortest path to a state each time we saw it.
- **Dynamic programming** - a table-driven method to solve problems by combining solutions to subproblems. (Eg. Viterbi algorithm)
- MED base cases:
  - $D[i,j]$  = The edit distance between  $X[1..i]$  and  $Y[1..j]$
  - $D[i,0] = i$  (requires  $i$  deletes)
  - $D[0,j] = j$  (requires  $j$  deletes)

## How to calculate MED?

$$D[i, j] = \min \begin{cases} D[i-1, j] + \text{del-cost}(\text{source}[i]) \\ D[i, j-1] + \text{ins-cost}(\text{target}[j]) \\ D[i-1, j-1] + \text{sub-cost}(\text{source}[i], \text{target}[j]) \end{cases}$$

### Levenshtein Distance

$$D[i, j] = \min \begin{cases} D[i-1, j] + 1 \\ D[i, j-1] + 1 \\ D[i-1, j-1] + \begin{cases} 2; & \text{if } \text{source}[i] \neq \text{target}[j] \\ 0; & \text{if } \text{source}[i] = \text{target}[j] \end{cases} \end{cases}$$



		0	1	2	3	4	5	6	7	8	9
Src\Tar		#	e	x	e	c	u	t	i	o	n
0	#	0	1	2	3	4	5	6	7	8	9
1	i	1	2	3	4	5	6	7	6	7	8
2	n	2	3	4	5	6	7	8	7	8	7
3	t	3	4	5	6	7	8	7	8	9	8
4	e	4	3	4	5	6	7	8	9	10	9
5	n	5	4	5	6	7	8	9	10	11	10
6	t	6	5	6	7	8	9	8	9	10	11
7	i	7	6	7	8	9	10	9	8	9	10
8	o	8	7	8	9	10	11	10	9	8	9
9	n	9	8	9	10	11	12	11	10	9	8

**Figure 2.18** Computation of minimum edit distance between *intention* and *execution* with the algorithm of Fig. 2.17, using Levenshtein distance with cost of 1 for insertions or deletions, 2 for substitutions.

	#	e	x	e	c	u	t	i	o	n
#	0	← 1	← 2	← 3	← 4	← 5	← 6	← 7	← 8	← 9
i	↑ 1	↖←↑ 2	↖←↑ 3	↖←↑ 4	↖←↑ 5	↖←↑ 6	↖←↑ 7	↖ 6	← 7	← 8
n	↑ 2	↖←↑ 3	↖←↑ 4	↖←↑ 5	↖←↑ 6	↖←↑ 7	↖←↑ 8	↑ 7	↖←↑ 8	↖ 7
t	↑ 3	↖←↑ 4	↖←↑ 5	↖←↑ 6	↖←↑ 7	↖←↑ 8	↖ 7	←↑ 8	↖←↑ 9	↑ 8
e	↑ 4	↖ 3	← 4	↖← 5	← 6	← 7	←↑ 8	↖←↑ 9	↖←↑ 10	↑ 9
n	↑ 5	↑ 4	↖←↑ 5	↖←↑ 6	↖←↑ 7	↖←↑ 8	↖←↑ 9	↖←↑ 10	↖←↑ 11	↖↑ 10
t	↑ 6	↑ 5	↖←↑ 6	↖←↑ 7	↖←↑ 8	↖←↑ 9	↖ 8	← 9	← 10	←↑ 11
i	↑ 7	↑ 6	↖←↑ 7	↖←↑ 8	↖←↑ 9	↖←↑ 10	↑ 9	↖ 8	← 9	← 10
o	↑ 8	↑ 7	↖←↑ 8	↖←↑ 9	↖←↑ 10	↖←↑ 11	↑ 10	↑ 9	↖ 8	← 9
n	↑ 9	↑ 8	↖←↑ 9	↖←↑ 10	↖←↑ 11	↖←↑ 12	↑ 11	↑ 10	↑ 9	↖ 8

At each step:

- If you follow ↖ (**diagonal**), you either substitute (if letters differ) or match (if same).
- If you follow ↑ (**top**), you've deleted a character from the source.
- If you follow ← (**left**), you've inserted a character into the source.

I N T E \* N T I O N

| | | | | | | | |

\* E X E C U T I O N

d s s i s

## When to Skip Preprocessing?

- **Context:** Tasks needing case sensitivity (e.g., NER).
- **Models:** Modern LLMs (BERT, GPT) handle raw text better.

## Impact on Model Performance

- Text preprocessing improves NLP system performance
- **Case Study:** Sentiment analysis accuracy improves by 15% after stop word removal.

**Key Takeaway:** Preprocessing tailors text for NLP tasks.

# Wider reading

Tokenization:

- Wordpiece: <https://huggingface.co/learn/llm-course/chapter6/6>
- Unigram: <https://huggingface.co/learn/llm-course/en/chapter6/7>

## References:

- Chapter 2: <https://web.stanford.edu/~jurafsky/slp3/ed3book.pdf>